

# **MATRIX MULTIPLICATION AND PERFORMANCE**

A Short Journey into High Performance Computing (HPC)

---

Kevin Waters

# MATRIX MULTIPLICATION AND PERFORMANCE

The Time I Accidentally Bested NumPy

---

Kevin Waters

- The code and talk for this presentation can be found [here](#).

A screenshot of a GitHub repository page. The repository name is "kwaters/Cleanup README". The commit history shows 14 commits from "kwaters" made 2 minutes ago. The commits are:

- data: First draft complete (15 minutes ago)
- include: First draft complete (15 minutes ago)
- lib: Add typst presentation, data, and python (2 hours ago)
- python: First draft complete (15 minutes ago)
- src: First draft complete (15 minutes ago)
- test: Add typst presentation, data, and python (2 hours ago)
- typsides: Cleanup README (2 minutes ago)
- CMakeLists.txt: Change to clang compiler (yesterday)
- README.md: Cleanup README (2 minutes ago)

The repository has 1 branch and 0 tags. The "About" section includes:

- Small exploration of Matrix Multiply
- Readme
- Activity
- 0 stars
- 0 watching
- 0 forks

The "Releases" section indicates "No releases published" and provides a link to "Create a new release". The "Packages" section indicates "No packages published" and provides a link to "Publish your first package". The "Languages" section shows the following distribution:

Language	Percentage
C	38.4%
Typst	35.4%
Python	23.6%
CMake	2.3%
TeX	0.3%

**Matrix Multiplication Exploration**

# CONTENTS


---

1. What is High Performance Computing (HPC)?
2. Matrix Multiplication
3. Matrix Multiplication in Python (Easy Part)
4. Matrix Multiplication in C (Hard Part)
5. More Discussion

# WHAT IS HIGH PERFORMANCE COMPUTING (HPC)?

---

- AI (Training/Inference)?
- Large-scale distributed memory computations?
- Distributed and scalable web services?
- Performance-aware programming:<sup>1</sup>
  - ▶ x86 aware?
  - ▶ Platform aware (CPU vs. GPU)?
  - ▶ Instruction set architecture (ISA) aware?
  - ▶ Cache-size aware?



*Raspberry Pi cluster<sup>2</sup>*

<sup>1</sup>Term taken from Casey Muratori <https://www.computerenhance.com/p/welcome-to-the-performance-aware>

<sup>2</sup><https://www.raspberrypi.com/tutorials/cluster-raspberry-pi-tutorial/>

# MATRIX MULTIPLICATION

---

- A few examples of applications (Where is linear algebra used today?):
  - **Finite element analysis** - Aerospace, automotive, material's properties...
  - **Electronic structure theory** - Density Functional Theory, Hartree-Fock++...
  - **Machine learning/data science** - Data analysis, pattern recognition, neural nets...
  - **Genetics** - genotype distribution
  - **Solving (partial) differential equations**... and much more

$$\mathbf{AB} = \mathbf{C}$$

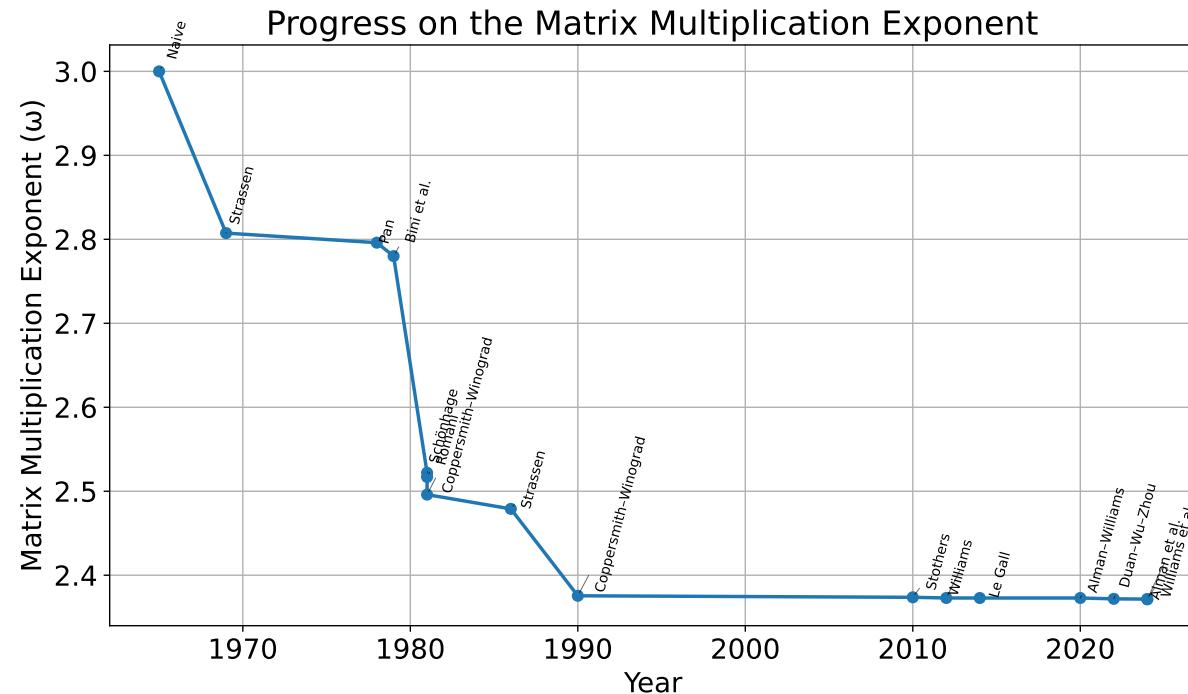
$$\begin{pmatrix} a_{00} & a_{01} & a_{02} & \dots & a_{0n} \\ a_{10} & a_{11} & a_{12} & \dots & a_{1n} \\ a_{20} & a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \vdots & \ddots & \\ a_{n0} & a_{n1} & a_{n3} & \dots & a_{nn} \end{pmatrix} * \begin{pmatrix} b_{00} & b_{01} & b_{02} & \dots & b_{0n} \\ b_{10} & b_{11} & b_{12} & \dots & b_{1n} \\ b_{20} & b_{21} & b_{22} & \dots & b_{2n} \\ \vdots & \vdots & \vdots & \ddots & \\ b_{n0} & b_{n1} & b_{n3} & \dots & b_{nn} \end{pmatrix} = \begin{pmatrix} c_{00} & c_{01} & c_{02} & \dots & c_{0n} \\ c_{10} & c_{11} & c_{12} & \dots & c_{1n} \\ c_{20} & c_{21} & c_{22} & \dots & c_{2n} \\ \vdots & \vdots & \vdots & \ddots & \\ c_{n0} & c_{n1} & c_{n3} & \dots & c_{nn} \end{pmatrix}$$

Where:

$$c_{ij} = \sum a_{ij} b_{kj}$$

# Computational Complexity

10 / 49



*Progress of Computation Complexity of Matrix Multiply<sup>3</sup>*

<sup>3</sup>Sourced from [https://en.wikipedia.org/wiki/Computational\\_complexity\\_of\\_matrix\\_multiplication](https://en.wikipedia.org/wiki/Computational_complexity_of_matrix_multiplication)

- Naive Matrix multiplication's computational complexity is  $\theta(n^3)$ 
  - ▶ Others scale better, but there are trade-offs
- What does that mean?
  - ▶ 2x2 results in 8 multiplication steps and 4 addition steps:

$$\begin{pmatrix} a_{00} & a_{01} \\ a_{10} & a_{11} \end{pmatrix} \begin{pmatrix} b_{00} & b_{01} \\ b_{10} & b_{11} \end{pmatrix} = \begin{pmatrix} a_{00}b_{00} + a_{01}b_{10} & a_{00}b_{01} + a_{01}b_{11} \\ a_{10}b_{00} + a_{11}b_{10} & a_{10}b_{01} + a_{11}b_{11} \end{pmatrix}$$

- ▶ For the general case the number of operations is given by the following:

$$2n^3 + n^2$$

- ▶  $2n^3$  multiply operations.
- ▶  $n^2$  addition operations.

# MATRIX MULTIPLICATION IN PYTHON (EASY PART)

---

## Python Benchmark

```
> python python_plain.py
1024x1024 matrix multiply...
Int      Time: 80.576309 seconds
Float    Time: 97.653700 seconds
```

- There are many things working against python, just-in-time-compilation (JIT), arbitrary size integers...

- NumPy is a highly-tuned library where typically expensive functions are implemented in compiled languages such as C/C++ or FORTRAN. (BLAS & LAPACK)

## Python (NumPy) Benchmark

```
> python python_numpy.py
1024x1024 matrix multiply...
Int32      Time: 1.895825 seconds
Int64      Time: 1.884659 seconds
Float       Time: 0.005361 seconds
Double      Time: 0.010751 seconds
```

- These results are interesting...

# MATRIX MULTIPLICATION IN C (HARD PART)

---

- All matrices are square.
- All matrices sizes are a power of 2.
- Matrices are populated with ints (32-bits) randomly distributed [-5,5].
- For benchmarking, one warm-up was followed by five trials (times are per trials).
- To count cycles, the following x86 instruction was used “rdtsc()” (Time Stamp Counter).
- Time stamps were called using `clock_gettime()`.
- Compile flags for all timed runs:
  - ▶ `-Wall -O3 -march=native -funroll-loops`
- Work per cycle was calculated using the following:

$$\frac{\text{Work Required}}{\text{Clock Cycle}} = \frac{3n^3 + n^2}{\# \text{cycles\_elapsed}}$$

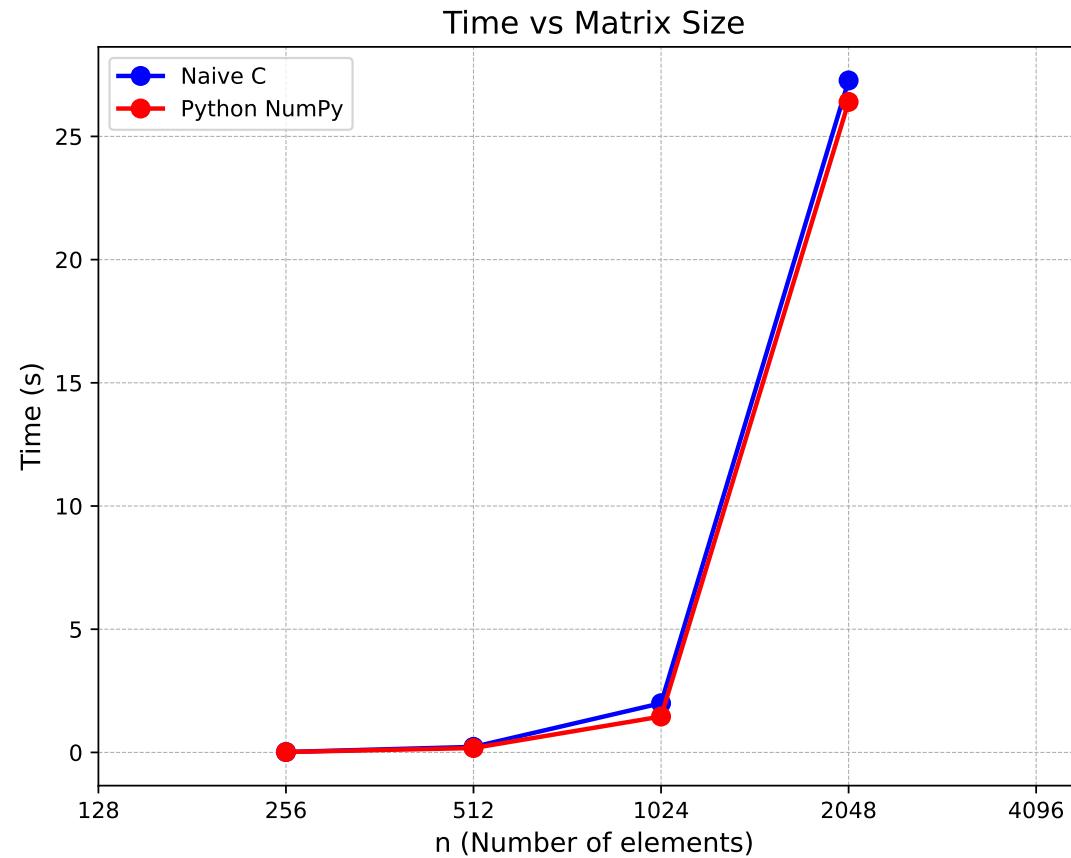
## Simple Matrix Multiply

```
for (int i = 0; i < n; i++) {  
    for (int j = 0; j < n; j++) {  
        for (int k = 0; k < n; k++) {  
            result[i * n + j] += matrix1[i * n + k] * matrix2[k * n + j];  
        }  
    }  
}
```

- Readable, simple, and slow.

# Naive Matrix Multiply

18 / 49




- Assuming 32-bit integers or 32-bit floats

<b>n</b>	<b>Elements (<math>n^2</math>)</b>	<b>Memory/Matrix (KiB)</b>	<b>Total Memory (KiB)</b>
64	4096	16	48
128	16384	64	192
256	65536	256	768
512	262144	1024	3072
1024	1048576	4096	12288
2048	4194304	16384	49152
4096	16777216	65536	196608


- Total memory is 3x the memory per matrix.

# Memory Footprint Revisited

20 / 49



*Memory requirement for three  $n \times n$  matrices.*



*8-core client Skylake topology (\$ lstopo)*

- One NUMA domain
- One 16 MiB shared L3 cache
- Individual 256 KiB L2 cache
- Individual 32 KiB instruction and data caches
- Two logical cores per physical core (8 physical, 16 logical)

# Latency Comparison Numbers (Everything is a Cache) (~2012)

23 / 49

L1 cache reference	0.5 ns			
Branch mispredict	5 ns			
L2 cache reference	7 ns			14x L1 cache
Mutex lock/unlock	25 ns			
Main memory reference	100 ns			20x L2 cache, 200x L1 cache
Send 1K bytes over 1 Gbps network	10,000 ns	10 us		
Read 4K randomly from SSD	150,000 ns	150 us		1GB/sec SSD
Read 1 MB sequentially from memory	250,000 ns	250 us		
Read 1 MB sequentially from SSD	1,000,000 ns	1,000 us	1 ms	1GB/sec SSD, 4X memory
Disk seek	10,000,000 ns	10,000 us	10 ms	10x datacenter roundtrip
Read 1 MB sequentially from disk	20,000,000 ns	20,000 us	20 ms	80x memory, 20X SSD

Latencies to generate intuition for the cost of an operation.<sup>4</sup>

<sup>4</sup>Originally by Peter Norvig: <http://norvig.com/21-days.html#answers>

## Intermediate Sum

```
> perf stat -e cycles,instructions,cache-references,cache-misses ./build/driver 2048 2048 5  
n, trials, req. memory (KiB), time/trial (s), work/cycle, Read Bandwidth (GiB/s)  
s2048, 5, 49152, 21.019992, 0.215482, 0.000297
```

Performance counter stats for './build/driver 2048 2048 5':

603,229,008,482	cycles:u	
156,197,169,903	instructions:u	# 0.26 insn per cycle
100,483,589,712	cache-references:u	
7,718,024,466	cache-misses:u	# 7.68% of all cache refs

- More targeted data can be found with PAPI (Performance Application Programming Interface).


## Intermediate Sum

```
for (int i = 0; i < n; i++) {  
    for (int j = 0; j < n; j++) {  
        int sum = 0;  
        for (int k = 0; k < n; k++) {  
            sum = matrix1[i * n + k] * matrix2[k * n + j] + sum;  
        }  
        result[i * n + j] = sum;  
    }  
}
```

- Allows for sum to stay in registers requiring less fetching from memory, a little bit faster.

# Intermediate Sum Matrix Multiply

26 / 49




## Loop Re-ordering Sum

```
for (int i = 0; i < n; i++) {  
    for (int k = 0; k < n; k++) {  
        for (int j = 0; j < n; j++) {  
            result[i * n + j] += matrix1[i * n + k] * matrix2[k * n + j];  
        }  
    }  
}
```

- Re-ordering the last two loops (j with k) enabling better caching behavior.

# Loop Re-ordering Matrix Multiply

28 / 49




## Blocking

```
for (int ii = 0; ii < n; ii+= BLOCK_SIZE) {
    for (int kk = 0; kk < n; kk+= BLOCK_SIZE) {
        for (int jj = 0; jj < n; jj+= BLOCK_SIZE) {
            int limit_i = ((ii + BLOCK_SIZE) < n) ? (ii + BLOCK_SIZE) : n;
            int limit_j = ((jj + BLOCK_SIZE) < n) ? (jj + BLOCK_SIZE) : n;
            int limit_k = ((kk + BLOCK_SIZE) < n) ? (kk + BLOCK_SIZE) : n;
            for (int i = ii; i < limit_i; ++i) {
                for (int k = kk; k < limit_k; ++k) {
                    int ki = i * n + k;
                    for (int j = jj; j < limit_j; j++) {
                        result[i * n + j] += matrix1[ki] * matrix2[k * n + j];
                    }
                }
            }
        }
    }
}
```

# Block Matrix Multiply

30 / 49




*GEMM tiling or BLOCK\_SIZE.<sup>5</sup>*

<sup>5</sup>Mathhes et. al. Tuning and Optimization for a Variety of Many-Core Architectures Without Changing a Single Line of Implementation Code Using the Alpaka Library (2017)

- This is where optimizations start becoming unpleasant as it is not hardware agnostic, however we are not using intrinsics yet!
- The BLOCK\_SIZE variable is a configure time constant, requiring the library to be reconfigured and recompiled.
- We will recompile until we find an optimal BLOCK\_SIZE value.

# Loop Re-ordering Matrix Multiply BLOCK\_SIZE

32 / 49




## BLOCK\_SIZE 512

```
> perf stat -e cycles,instructions,cache-references,cache-misses ./build/driver 2048 2048 5  
2048, 5, 49152, 0.792592, 5.714694, 0.007886
```

```
Performance counter stats for './build/driver 2048 2048 5':
```

22,990,040,828	cycles:u	
26,574,845,684	instructions:u	# 1.16 insn per cycle
7,550,989,191	cache-references:u	
121,521,554	cache-misses:u	# 1.61% of all cache refs

- Naive L3 cache misses was 7.68%, PAPI would give better resolution. There is still better cache performance possible.



## Checking the Hardware

```
> for i in Architecture "CPU(s):" "Model name" Thread Socket "NUMA node(s)"; do  
lscpu | grep "$i" | grep -v "node0"; done
```

Architecture:	x86_64
CPU(s):	16
Model name:	Intel(R) Core(TM) i7-10700KF CPU @ 3.80GHz
Thread(s) per core:	2
Socket(s):	1
NUMA node(s):	1

- x86<sup>6</sup> SIMD<sup>7</sup> extensions:
  - **SSE (Streaming SIMD Extensions)** - 128-bit floating point registers
  - **SSE2** - 128-bit doubles and integer registers
  - **AVX (Advanced Vector Extensions)** - 256-bit floating/double point registers
  - **AVX2** - 256-bit integer SSE instructions
  - **AVX512** - 512 bit registers!
- **FMA(Fused Multiply-Add)** - Exists in AVX, AVX512, but only for floats and doubles.

$$\text{result} = a * b + c$$

---

<sup>6</sup>ARM has different names for everything

<sup>7</sup>Single Instruction, Multiple Data

## Checking the Architecture and ISA

```
> cat /sys/devices/cpu/caps/pmu_name  
skylake  
  
> for isa in sse sse2 avx avx2 avx512 fma; do grep -q "$isa" /proc/cpuinfo && echo "$isa 1" ||  
echo "$isa 0"; done  
sse 1  
sse2 1  
avx 1  
avx2 1  
avx512 0  
fma 1
```

## Checking for SIMD (AVX2)


```
10b0: c4 c2 7d 40 4c 8a a0 vpmulld -0x60(%r10,%rcx,4),%ymm0,%ymm1
10b7: c4 c2 7d 40 54 8a c0 vpmulld -0x40(%r10,%rcx,4),%ymm0,%ymm2
10be: c4 c2 7d 40 5c 8a e0 vpmulld -0x20(%r10,%rcx,4),%ymm0,%ymm3
10c5: c4 c2 7d 40 24 8a vpmulld (%r10,%rcx,4),%ymm0,%ymm4
10cb: c4 c1 75 fe 4c 88 a0 vpaddd -0x60(%r8,%rcx,4),%ymm1,%ymm1
10d2: c4 c1 6d fe 54 88 c0 vpaddd -0x40(%r8,%rcx,4),%ymm2,%ymm2
10d9: c4 c1 65 fe 5c 88 e0 vpaddd -0x20(%r8,%rcx,4),%ymm3,%ymm3
10e0: c4 c1 5d fe 24 88 vpaddd (%r8,%rcx,4),%ymm4,%ymm4
10e6: c4 c1 7e 7f 4c 88 a0 vmovdqu %ymm1,-0x60(%r8,%rcx,4)
10ed: c4 c1 7e 7f 54 88 c0 vmovdqu %ymm2,-0x40(%r8,%rcx,4)
10f4: c4 c1 7e 7f 5c 88 e0 vmovdqu %ymm3,-0x20(%r8,%rcx,4)
10fb: c4 c1 7e 7f 24 88 vmovdqu %ymm4,(%r8,%rcx,4)
```

- Excerpt from multiplication library (appears some pipelining is going on!)<sup>8</sup>.

<sup>8</sup>More details can be found in the Intel Intrinsics Guide

# What is Good Performance?

39 / 49



AVX2 Throughput:

$$256 \frac{\text{bit}}{\text{cycle}} * \frac{1 \text{ byte}}{8 \text{ bit}} \frac{1 \text{ int}}{4 \text{ byte}} = 8 \frac{\text{int}}{\text{cycle}}$$

- This is a complicated question (multiply + addition):

$$c_{ij} = \sum a_{ij} b_{kj}$$

---

```
_m256i _mm256_add_epi32 (_m256i a, _m256i b)
```

#### Synopsis

```
_m256i _mm256_add_epi32 (_m256i a, _m256i b)
#include <immintrin.h>
Instruction: vpaddd ymm, ymm, ymm
CPUID Flags: AVX2
```

#### Description

Add packed 32-bit integers in **a** and **b**, and store the results in **dst**.

#### Operation

```
FOR j := 0 to 7
    i := j*32
    dst[i+31:i] := a[i+31:i] + b[i+31:i]
ENDFOR
dst[MAX:256] := 0
```

#### Latency and Throughput

Architecture	Latency	Throughput (CPI)
Alderlake	1	0.333333333
Icelake Intel Core	1	0.33
Icelake Xeon	1	0.33
Sapphire Rapids	1	0.333333333
Skylake	1	0.33

---

```
_m256i _mm256_mullo_epi32 (_m256i a, _m256i b)
```

#### Synopsis

```
_m256i _mm256_mullo_epi32 (_m256i a, _m256i b)
#include <immintrin.h>
Instruction: vpmulld ymm, ymm, ymm
CPUID Flags: AVX2
```

#### Description

Multiply the packed signed 32-bit integers in **a** and **b**, producing intermediate results in **tmp**.

#### Operation

```
FOR j := 0 to 7
    i := j*32
    tmp[63:0] := a[i+31:i] * b[i+31:i]
    dst[i+31:i] := tmp[31:0]
ENDFOR
dst[MAX:256] := 0
```

#### Latency and Throughput


Architecture	Latency	Throughput (CPI)
Alderlake	10	1
Icelake Intel Core	-	1
Icelake Xeon	10	1
Sapphire Rapids	-	1
Skylake	10	0.66

- Taken from the Intel Intrinsics Guide<sup>9</sup>

<sup>9</sup><https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html>


# Max Performance Estimation (x86 Is Hard)

42 / 49



*Cut out from the Skylake micro-architecture<sup>10</sup>*

<sup>10</sup>[https://en.wikichip.org/wiki/intel/microarchitectures/skylake\\_%28client%29](https://en.wikichip.org/wiki/intel/microarchitectures/skylake_%28client%29)



<b>Matrix size (n)</b>	<b>Runtime (s)</b>		
	<b>Pure Python</b>	<b>Python (NumPy)</b>	<b>Blocked C</b>
256	0.24	0.02	0.001
512	1.79	0.18	0.01
1024	15.32	1.46	0.10
2048		26.40	0.81
4096			6.55

- More optimization is possible for C and NumPy. Beating NumPy in this case is easy.

- Ask yourself, does it need to be further optimized?
  - Has someone else done it?
- Pack matrix B into contiguous blocks of memory.
  - Row major matrices, it is being accessed as a column.
  - May assist in better pipelining
- Continue additional block techniques to better utilize the L1 and L2 caches.
- Start using more cores, making sure to avoid race conditions and atomics (if possible)
- On the surface, ints will be 1/2 as fast as floats (due to FMA)
- Add generics to enable multiple data types for a more robust library.

## MORE DISCUSSION

---

- Libraries exists with hyper optimized floating-point matrix operations standardized by BLAS<sup>11</sup>.
  - **ATLAS** - Automatically Tuned Linear Algebra Software
  - **OpenBLAS** - open-source CPU based BLAS
  - **rocBLAS** - AMD's GPUs version via ROCM
  - and many more<sup>12</sup>
- Sparsity may drive to different algorithms.
- If working with integers you may have to write your own kernels.
- If working with Boolean matrices they allow for new algorithms using look-up tables<sup>13</sup>.

---

<sup>11</sup>Basic Linear Algebra Subprograms

<sup>12</sup>[https://en.wikipedia.org/wiki/Basic\\_Linear\\_Algebra\\_Subprograms#Implementations](https://en.wikipedia.org/wiki/Basic_Linear_Algebra_Subprograms#Implementations)

<sup>13</sup>Method of Four Russians

- **OpenMP/PThreads**
  - Using all cores on a socket/node
- **Simultaneous Multithreading (SMT/HyperThreading)**
  - Should it be used for this application?
- **Non-Uniform Memory Access (NUMA)**
  - Even more levels to the memory subsystem
  - AMD's Core Complex (CCX) have made this harder
- **MPI/SHMEM**
  - Inter-node communication using remote direct memory access (RDMA)

# Questions?